




CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Distributed TensorFlow: A performance evaluation

End-of-internship Seminar

Emanuele Bugliarello

emanuele.bugliarello@gmail.com

September 6, 2017

<https://github.com/e-bug/distributed-tensorflow-benchmarks> 

Introduction

What is TensorFlow?

Google's open-source software library for Machine Learning

- Best-supported client language: Python
- Experimental interfaces for: C++, Java and Go

Why TensorFlow?

- Portable & flexible → popular in industries and in research communities
- Most CSCS clients choose TensorFlow as their Deep Learning library

Why distributed training?

Training a neural network can take an impractically long time on a single machine (even with a GPU)

Results

On 64 GPUs: ~80% scalability efficiency in Piz Daint & almost 90% in 8 8-GPU nodes

ToC

- Introduction
- Distributed training in TensorFlow
- Benchmarks
- Conclusion and Future Work

Distributed training in TensorFlow (1)

TensorFlow is based on data flow graphs

- Nodes represent mathematical operations
- Tensors move across the edges between nodes

Writing a TensorFlow application

1. Build computation graph
2. Run instances of that graph

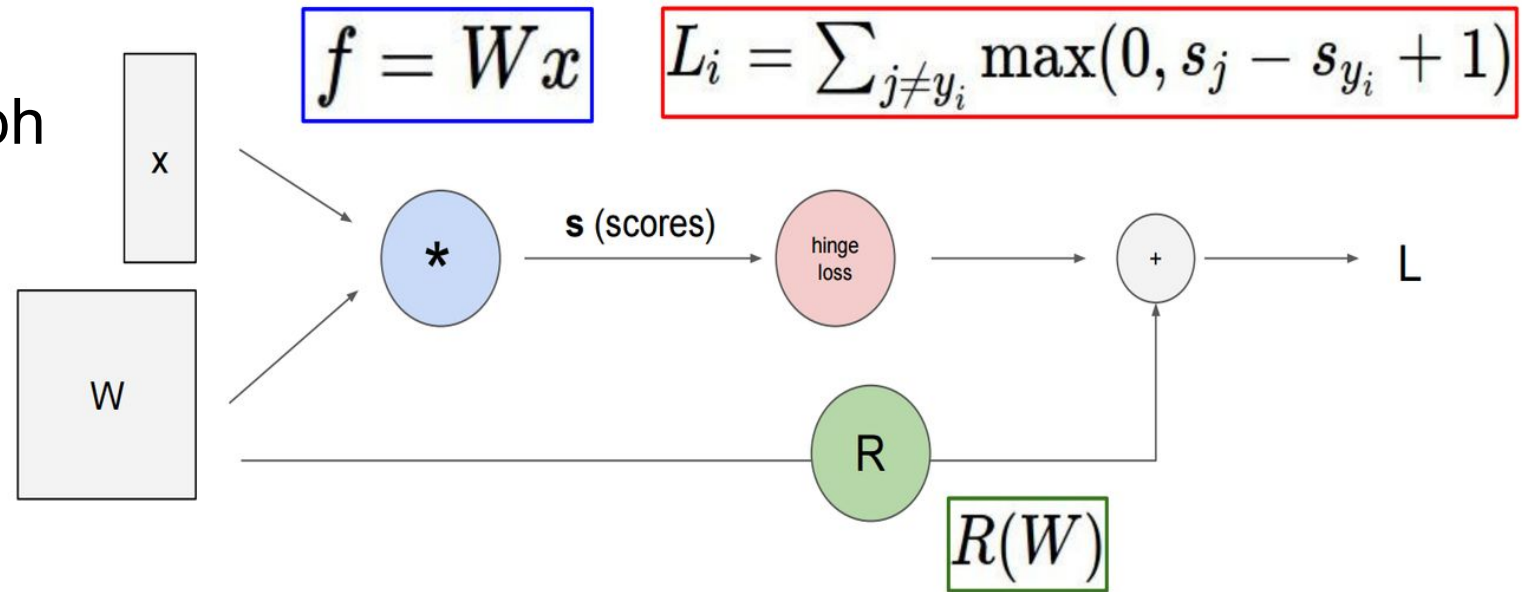
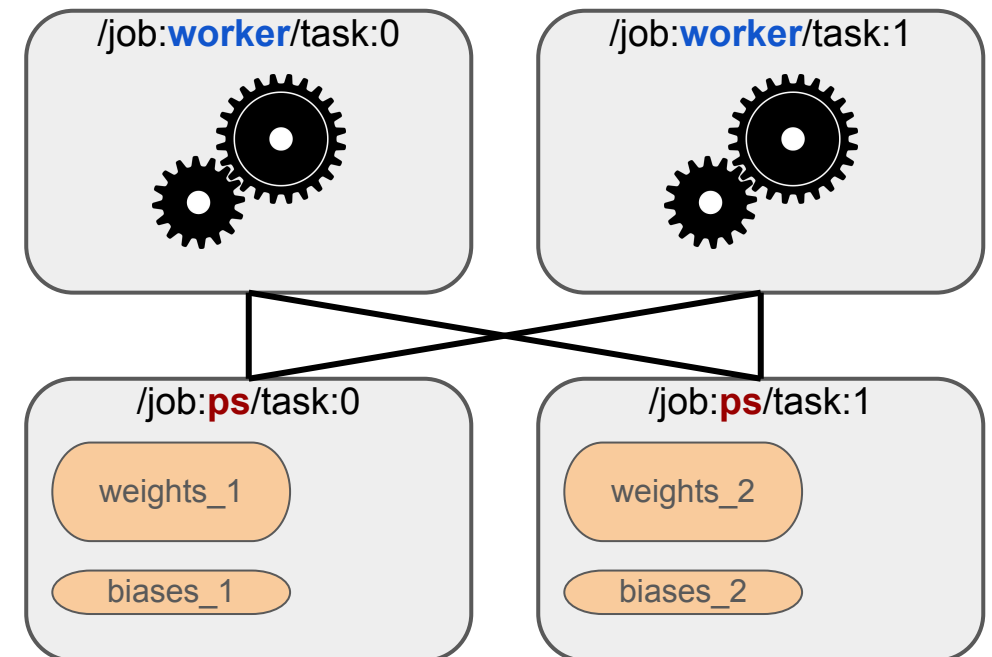


Figure 1: Computational graph for regularized Multiclass SVM loss ([CS231N, Stanford University](#))



Distributed training in TensorFlow (2)

- Split the training of a neural network across multiple nodes
- Most common approach: data parallelism
 - Each node has an instance of the model and reads different training samples
 - Also known as “between-graph replication” in TensorFlow
- Processes can either be:
 - Worker
 - Runs the model
 - Sends its local gradients to the PSs
 - Receives updated variables back
 - Parameter Server (PS)
 - Hosts trainable variables
 - Updates them with values sent by the Workers
- PSs sum gradients to merge in one step what each Worker has learned to reduce the loss



Distributed training in TensorFlow (3)

- Workers need to send their updates to the correct Parameter Servers
 - Use TensorFlow's `replica_device_setter` for a deterministic variable allocation
- Parameter Servers and Workers may coexist on the same machine
 - Recommended when Workers run on GPUs
 - Reduce the number of nodes
 - Minimize network communications



Distributed training in TensorFlow (4)

- Define cluster of nodes and the role of each of them (PS/Worker)
- The following code snippet (<https://clindatsci.com/blog/2017/5/31/distributed-tensorflow>) would be executed on each machine in the cluster, but with different arguments

```
import sys
import tensorflow as tf

# Specify the cluster's architecture
cluster = tf.train.ClusterSpec(
    {'ps': ['192.168.1.1:1111'],
     'worker': ['192.168.1.2:1111', '192.168.1.3:1111']})

# Parse command-line to specify machine
job_type = sys.argv[1] # job type: "worker" or "ps"
task_idx = sys.argv[2] # index job in the worker or ps list
                        # as defined in the ClusterSpec

# Create TensorFlow Server.
# This is how the machines communicate.
server = tf.train.Server(cluster, job_name=job_type,
                        task_index=task_idx)
```

```
# Parameter server is updated by remote clients.
# Will not proceed beyond this if statement.
if job_type == 'ps':
    server.join()
else:
    # Workers only
    with tf.device(tf.train.replica_device_setter(
        worker_device='/job:worker/task:'+task_idx,
        cluster=cluster)):
        # Build your model here
        # as if you only were using a single machine

    with tf.Session(server.target):
        # Train your model here
```

Distributed training in TensorFlow (5)

Running distributed TensorFlow on Piz Daint

- Write a Python script (**TF_SCRIPT**) accepting **job_name**, **task_index**, **ps_hosts** and **worker_hosts** TensorFlow flags
- Write a Bash script like the following one; **run_dist_tf_daint.sh** will specify the cluster from allocated nodes

```
#!/bin/bash
```

```
#SBATCH --job-name=distributed_tf
#SBATCH --time=00:12:00
#SBATCH --nodes=8
#SBATCH --constraint=gpu
#SBATCH --output=distributed_tf.%j.log
```

```
# Arguments:
```

```
# $1: TF_NUM_PS: number of parameter servers
```

```
# $2: TF_NUM_WORKER: number of workers
```

```
# load modules
```

```
module load daint-gpu
```

```
module load TensorFlow
```

```
# set TensorFlow script parameters
```

```
export TF_SCRIPT="$HOME/project_dir/project_script.py"
```

```
export TF_FLAGS="--num_gpus=1 --batch_size=64  
--num_batches=4 --data_format=NCHW"
```

```
# set TensorFlow distributed parameters
```

```
export TF_NUM_PS=$1 # 1
```

```
export TF_NUM_WORKERS=$2 # $SLURM_JOB_NUM_NODES
```

```
# export TF_WORKER_PER_NODE=1
```

```
# export TF_PS_PER_NODE=1
```

```
# export TF_PS_IN_WORKER=true
```

```
# run distributed TensorFlow
```

```
DIST_TF_LAUNCHER_DIR=$SCRATCH/run_dist_tf_daint_dir
```

```
cd $DIST_TF_LAUNCHER_DIR
```

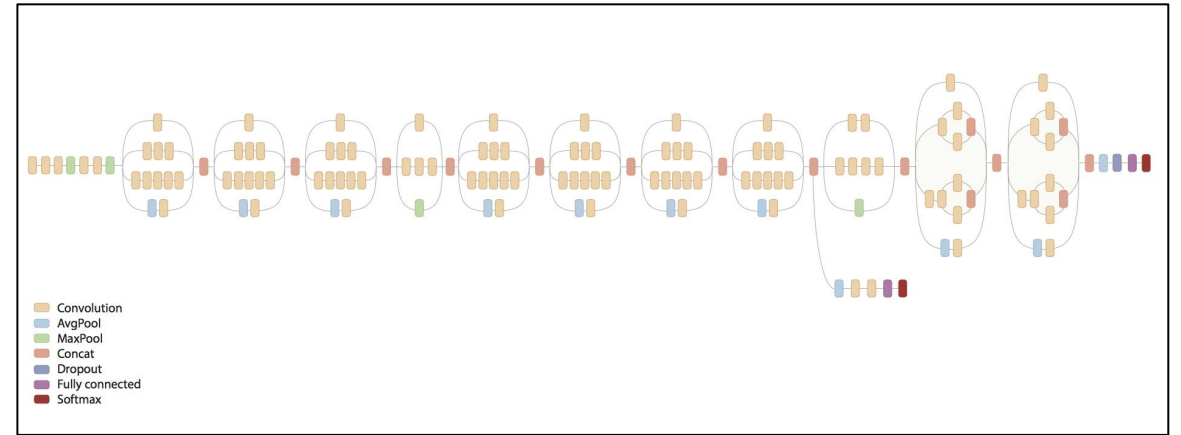
```
./run_dist_tf_daint.sh
```


ToC

- Introduction
- Distributed training in TensorFlow
- **Benchmarks**
- Conclusion and Future Work

Benchmarks (1)

- Model
 - InceptionV3
 - Neural Network for 1000-class image classification (ImageNet competition)
 - Optimized code for benchmarks available from Google
- Data set
 - ImageNet: 1,280,000 images (144 GB)
- TensorFlow 1.1.0
- Performance metric
 - Number of trained images per second



Benchmarks (2)


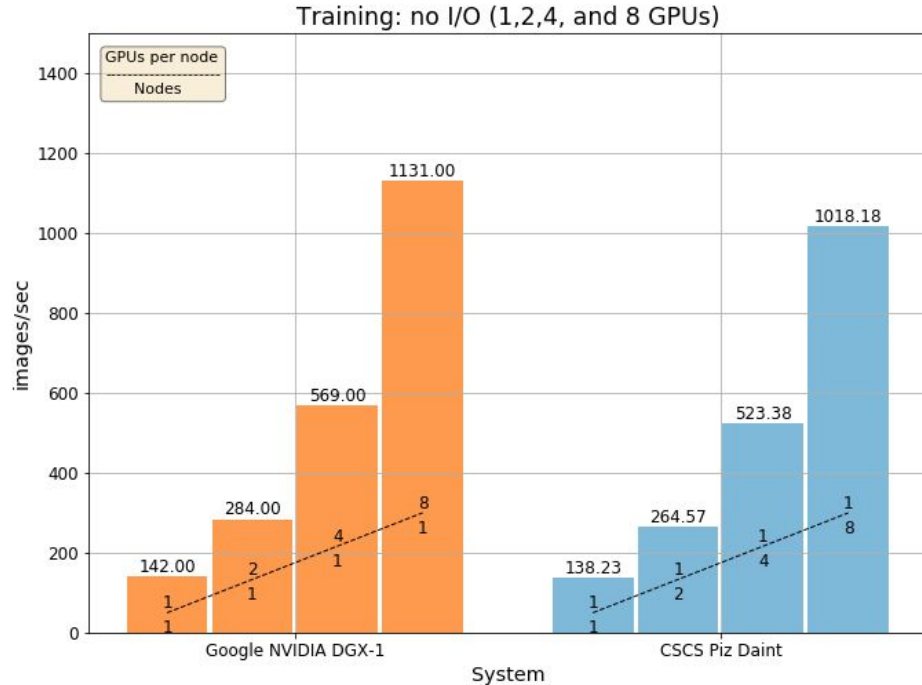
- Methodology
 - For each configuration of number of Workers and number of nodes
 - Run with different number of Parameter Servers on synthetic data (no I/O access)
 - Report best setting of number of Workers and number of PSs
 - Run best setting on real data (with I/O access)
 - Results repeatability
 - Run each test 5 times and average times together (Google's approach)
 - Compare results with Google's
 - Limit impact on Piz Daint (200+ tests)
 - For each test
 - 10 warmup steps
 - Next 100 steps are averaged

Benchmarks (3)

- Systems
 - Piz Daint (NVIDIA Tesla P100 - 1 GPU per node)
 - Amazon EC2 instances
 - p2.xlarge (NVIDIA Tesla K80 - 1 GPU per node)
 - p2.8xlarge (NVIDIA Tesla K80 - 8 GPUs per node)
- Benchmarks from Google available at <https://www.tensorflow.org/performance/benchmarks>
 - Google's systems
 - NVIDIA DGX-1 (NVIDIA Tesla P100 - 8 GPUs per node)
 - Amazon p2.8xlarge (NVIDIA Tesla K80 - 8 GPUs per node)

Benchmarks (4)

NVIDIA Tesla P100 - synthetic data (no I/O) - up to 8 GPUs




Scalability efficiency

- 99.56% on 8 GPUs in NVIDIA DGX-1
- 92.07% on 8 GPUs in Piz Daint
- 8 nodes in Piz Daint have similar performance as an NVIDIA DGX-1

Benchmarks (5)

NVIDIA Tesla K80 - synthetic data (no I/O) - up to 8 GPUs


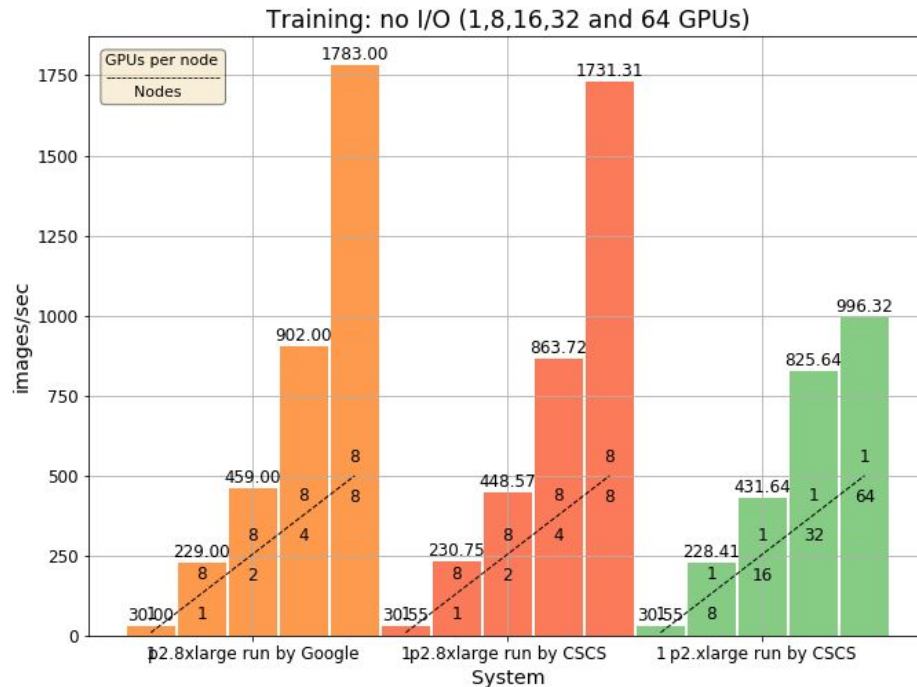


Scalability efficiency

- 94.58% and 94.44% on 8 GPUs in p2.8xlarge
- 93.45% on 8 GPUs in p2.xlarge
- Up to 8 GPUs, compute bound application

Benchmarks (6)

NVIDIA Tesla K80 - synthetic data (no I/O) - up to 64 GPUs




Scalability efficiency

- 92.86% and 88.55% on 64 GPUs in p2.8xlarge
- 50.96% on 64 GPUs in p2.xlarge
- Intuition: inter-node network capacity reached with 64 GPUs in p2.xlarge

Benchmarks (7)

Piz Daint (NVIDIA Tesla P100) - synthetic and real data - up to 128 GPUs

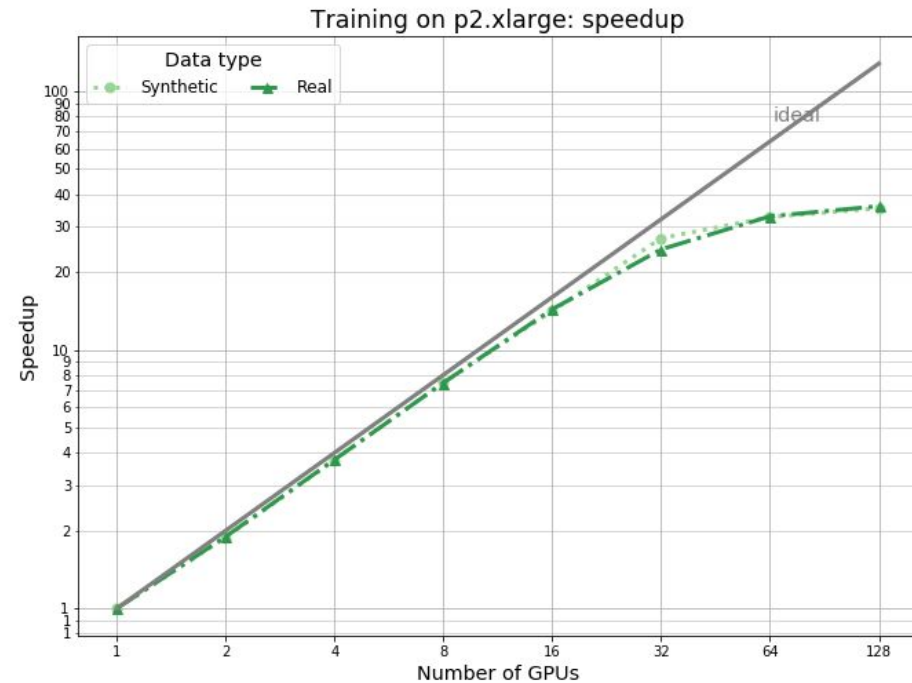
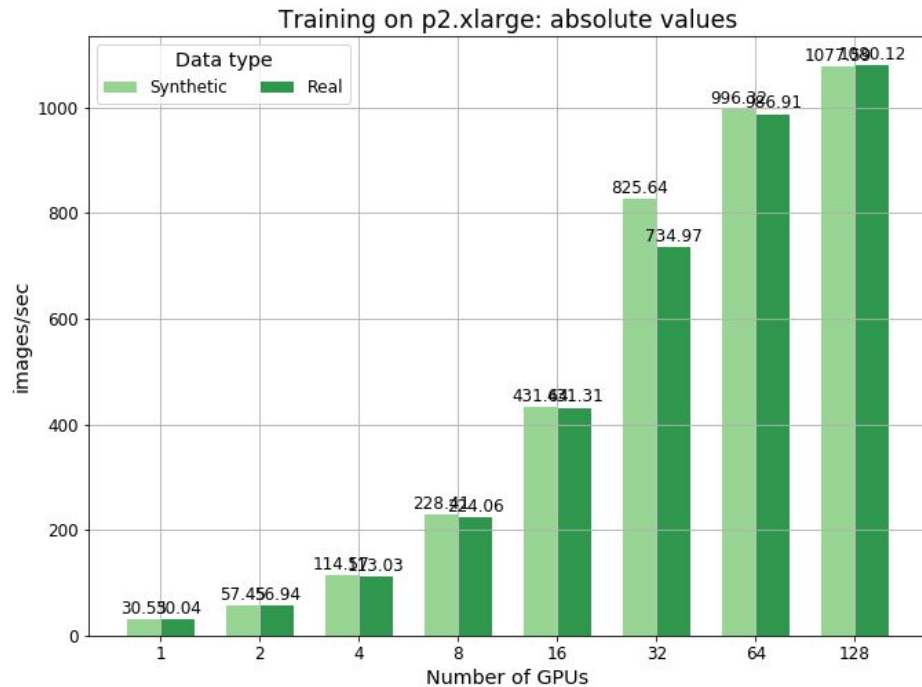


Scalability efficiency

- 80.46% (synthetic) and 72.39% (real) on 64 GPUs
- 52.11% (synthetic) and 51.63% (real) on 128 GPUs
- Intuition: inter-node network capacity reached with 128 nodes

Benchmarks (8)

p2.xlarge (NVIDIA Tesla K80) - synthetic and real data - up to 128 GPUs




Scalability efficiency (local SSD on each node)

- 50.96% (synthetic) and 51.33% (real) on 64 GPUs
- 27.56% (synthetic) and 28.09% (real) on 128 GPUs
- Intuition: inter-node network capacity reached with 64 nodes

Benchmarks (9)

p2.8xlarge (8 * NVIDIA Tesla K80) - synthetic and real data - up to 128 GPUs

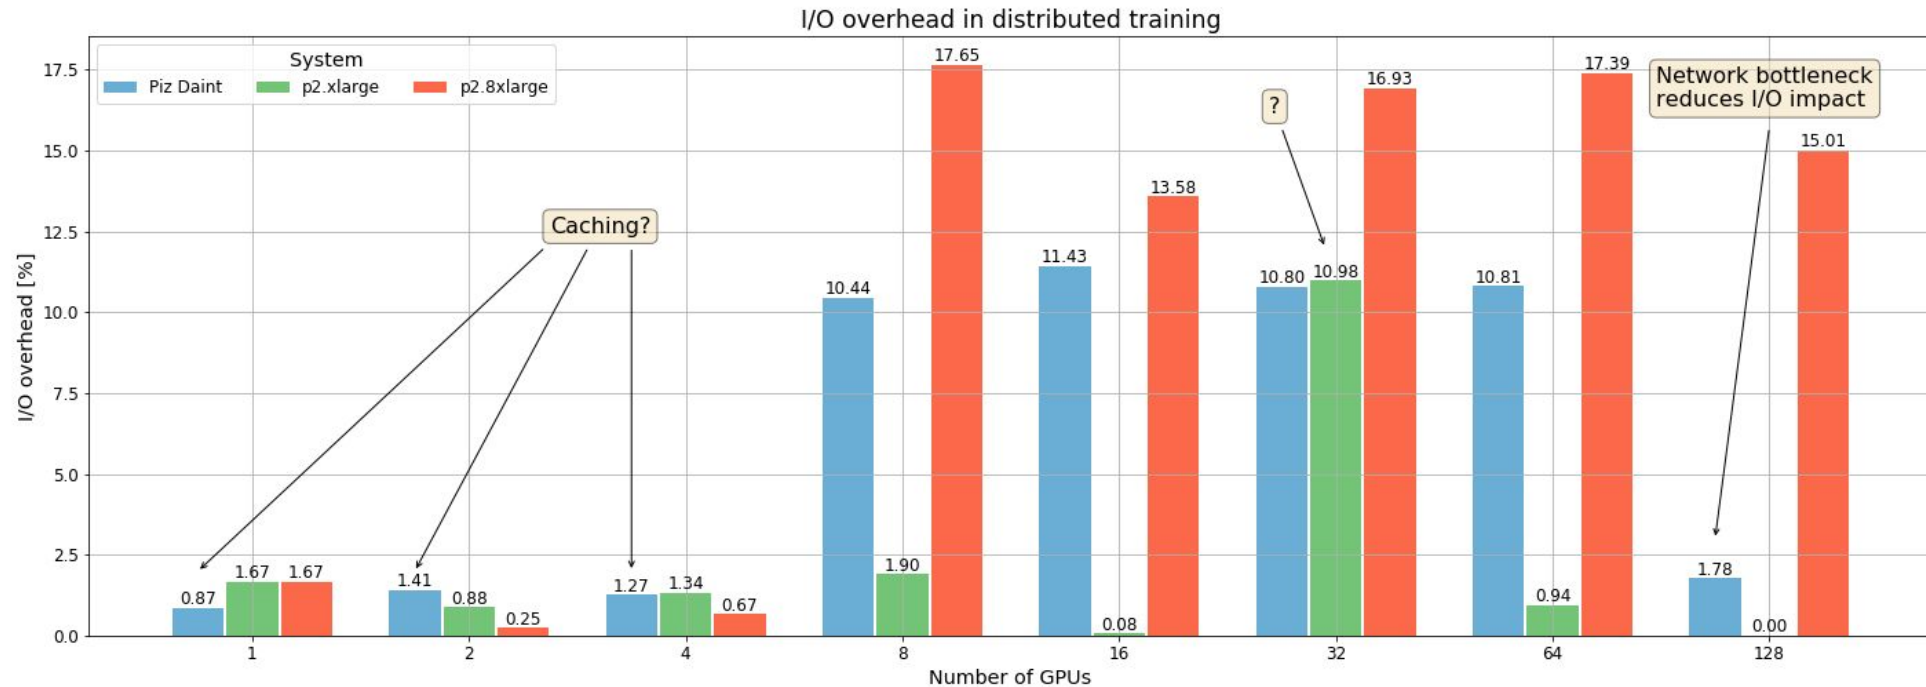


Scalability efficiency (local SSD on each node)

- 88.55% (synthetic) and 74.39% (real) on 64 GPUs
- 85.24% (synthetic) and 73.67% (real) on 128 GPUs
- Intuition: inter-node network capacity not reached (only 16 nodes for 128 GPUs)

Benchmarks (9)

I/O overhead



- ~17% on p2.8xlarge when 8 GPUs per node are used
- ~1% on p2.xlarge
- ~11% on Piz Daint when 8 to 64 GPUs used, ~1.5% otherwise

ToC

- Introduction
- Distributed training in TensorFlow
- Benchmarks
- Conclusion and Future Work

Conclusion

- 8 nodes in Piz Daint have similar performance to 1 NVIDIA DGX-1
- Scalability for InceptionV3 in TensorFlow
 - On Piz Daint
 - Supposedly inter-node bandwidth capacity reached after 64 nodes
 - I/O cost ~11%
 - On a multi-GPU system
 - Inter-node traffic algorithmically reduced by the number of GPUs per node (interconnect seems to have no real impact)
 - Using local SSDs and 8 GPUs per node adds a constant ~17% I/O overhead (PCIe traffic)
 - No benchmarks available for multiple NVIDIA DGX-1
- ⇒ Estimation according to the examined use case: Similar performance between 64 nodes on Piz Daint and 8 NVIDIA DGX-1 connected by a reasonable inter-node network

Future Work

- Investigate impact of training accuracy in distributed setting (preliminary results)
- Profile TensorFlow communication patterns
- Analyze influence of number of PSs for single- and multi-GPU systems

Conclusion

Thank you

- 8 nodes in Piz Daint have similar performance to 1 NVIDIA DGX-1
- Scalability for InceptionV3 in TensorFlow
 - On Piz Daint
 - Supposedly inter-node bandwidth capacity reached after 64 nodes
 - I/O cost ~11%
 - On a multi-GPU system
 - Inter-node traffic algorithmically reduced by the number of GPUs per node (interconnect seems to have no real impact)
 - Using local SSDs and 8 GPUs per node adds a constant ~17% I/O overhead (PCIe traffic)
 - No benchmarks available for multiple NVIDIA DGX-1
- ⇒ Estimation according to the examined use case: Similar performance between 64 nodes on Piz Daint and 8 NVIDIA DGX-1 connected by a reasonable inter-node network

Future Work

- Investigate impact of training accuracy in distributed setting (preliminary results)
- Profile TensorFlow communication patterns
- Analyze influence of number of PSs for single- and multi-GPU systems

Backup slides

TensorFlow overview (1)

TensorFlow is based on data flow graphs

- Nodes represent mathematical operations
- Tensors move across the edges between nodes

Writing a TensorFlow application

1. Build computation graph
2. Run instances of that graph




Figure 1: Computational graph for regularized Multiclass SVM loss ([CS231N, Stanford University](#))

TensorFlow overview (2)

Example: Linear Regression in TensorFlow




Figure 2: Computational graph for Linear Regression with squared loss

```
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
```

```
# ===== #
#                               #
#                               #
#                               #
# Generate some data as y=3*x + noise
N_SAMPLES = 10
x_in = np.arange(N_SAMPLES)
y_in = 3*x_in + np.random.randn(N_SAMPLES)
data = list(zip(x_in, y_in))
```

TensorFlow overview (3)

Example: Linear Regression in TensorFlow




Figure 2: Computational graph for Linear Regression with squared loss

```
# ===== #  
#                               #  
# BUILD GRAPH                  #  
# ===== #
```

```
simple_graph = tf.Graph()  
with simple_graph.as_default():  
    # Generate placeholders for input x and output y  
    x = tf.placeholder(tf.float32, name='x')  
    y = tf.placeholder(tf.float32, name='y')  
  
    # Create weight and bias, initialized to 0  
    w = tf.Variable(0.0, name='weight')  
    b = tf.Variable(0.0, name='bias')  
  
    # Build model to predict y  
    y_predicted = x * w + b  
  
    # Use the square error as the loss function  
    loss = tf.square(y - y_predicted, name='loss')  
  
    # Use gradient descent to minimize loss  
    optimizer = tf.train.GradientDescentOptimizer(0.001)  
    train = optimizer.minimize(loss)
```

TensorFlow overview (4)

Example: Linear Regression in TensorFlow

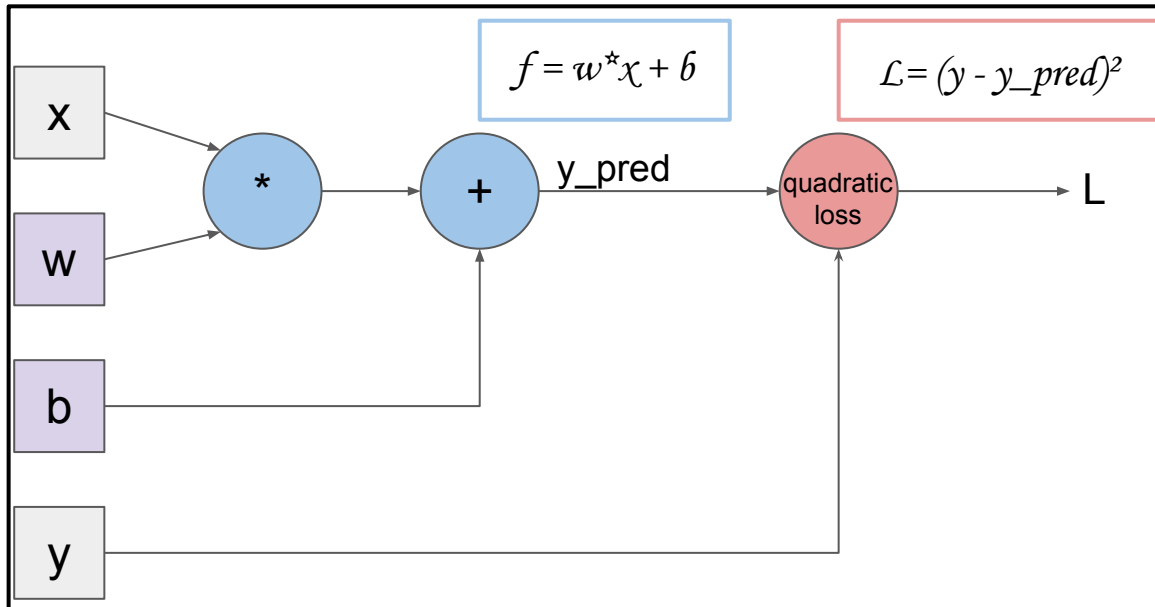


Figure 2: Computational graph for Linear Regression with squared loss

```
# ===== #  
#                               #  
# EXECUTE GRAPH                 #  
# ===== #  
# Run training for N_EPOCHS epochs  
N_EPOCHS = 5  
with tf.Session(graph=simple_graph) as sess:  
    # Initialize the necessary variables (w and b here)  
    sess.run(tf.global_variables_initializer())  
  
    # Train the model  
    for i in range(N_EPOCHS):  
        total_loss = 0  
        for x_, y_ in data:  
            # Session runs train operation and fetches values of loss  
            _, l_value = sess.run([train, loss], feed_dict={x: x_, y: y_})  
            total_loss += l_value  
            print('Epoch {0}: {1}'.format(i, total_loss/N_SAMPLES))  
  
    # Output final values of w and b  
    w_value, b_value = sess.run([w, b])
```

TensorFlow overview (5)

Example: Linear Regression in TensorFlow




Figure 2: Computational graph for Linear Regression with squared loss

```
# ===== #  
# PLOT RESULTS #  
# ===== #  
print(w_value, b_value) # 2.89, 0.45  
plt.plot(x_in, y_in, 'bo', label='Real data')  
plt.plot(x_in, x_in*w_value + b_value, 'orange',  
         label='Predicted data')  
plt.ylabel('y');plt.xlabel('x')  
plt.title('Linear Regression')  
plt.legend();plt.grid()  
plt.show()
```

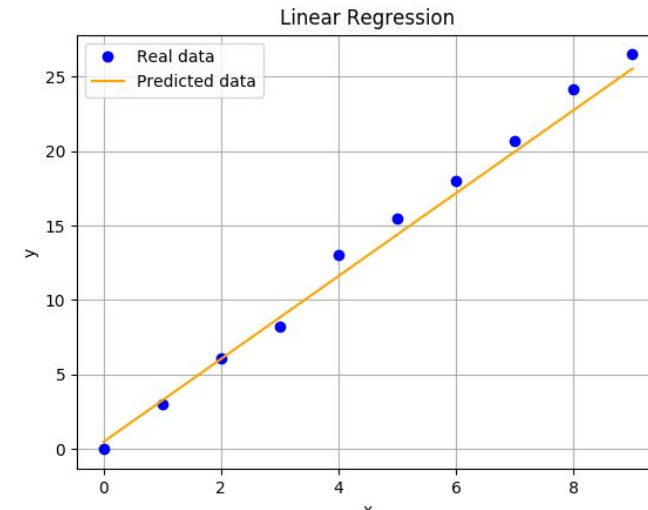
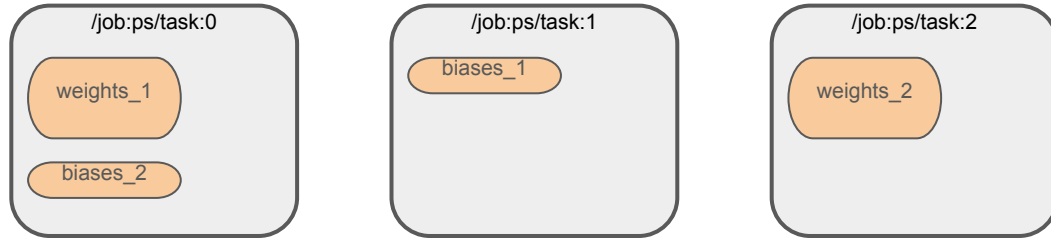



Figure 3: Learned linear model

Distributed training in TensorFlow (5)

Round-robin variables



Greedy load balancing variables



replica_device_setter provides two load balancing strategies

- Round-robin (default)
- Greedy load balancing

```
greedy = tf.contrib.training.GreedyLoadBalancingStrategy(...)
with tf.device(
    tf.train.replica_device_setter(ps_tasks=3,
                                   ps_strategy=greedy)):
    weights_1 = tf.get_variable('weights_1', [784, 100])
    biases_1 = tf.get_variable('biases_1', [100])
    weights_2 = tf.get_variable('weights_2', [100, 10])
    biases_2 = tf.get_variable('biases_2', [10])
```